

Trajectories Pipeline User Guide

Via Technology Ltd.

Trajectories Pipeline User Guide

© 2018 Via Technology Ltd

Phone: +44 1202 708476

Table of Contents

Revision History.....	4	4.6.3 Resources.....	18
References.....	5	4.7 Airport Intersections.....	18
Glossary.....	5	4.7.1 Input.....	18
1 Introduction.....	6	4.7.2 Output.....	18
2 System Setup.....	8	4.7.3 Resources.....	18
2.1 GCP Project.....	8	4.8 User Airspace Intersections.....	19
2.2 Google Bucket Storage.....	8	4.8.1 Input.....	19
2.3 Kubernetes Cluster.....	8	4.8.2 Output.....	19
2.4 Pipeline Container Image.....	9	4.8.3 Resources.....	19
2.4.1 Build Base Image.....	9	4.9 Import APDS Data.....	20
2.4.2 Build Python-only Image.....	9	4.9.1 Input.....	20
3 Data Organisation.....	10	4.9.2 Output.....	20
3.1 Bucket Storage Browser.....	11	4.9.3 Resources.....	20
3.2 gsutil.....	11	4.10 Merge APDS Data.....	20
4 Pipeline Processes.....	12	4.10.1 Input.....	21
4.1 Importing Daily Data.....	12	4.10.2 Output.....	21
4.1.1 Input.....	13	4.10.3 Resources.....	21
4.1.2 Output.....	13	5 Managing Jobs.....	22
4.1.3 Resources.....	13	5.1 Processing Order.....	22
4.2 Overnight Matching.....	14	5.2 Compute Resources.....	22
4.2.1 Inputs.....	14	5.2.1 High Memory Processes.....	23
4.2.2 Output.....	14	5.2.2 Low Memory Processes.....	24
4.2.3 Resources.....	14	5.2.3 Airspace Database.....	24
4.3 Overnight Merging.....	15	5.3 Running Processes.....	24
4.3.1 Input.....	15	5.3.1 Provisioning Compute Resources	24
4.3.2 Output.....	15	5.4 Running Jobs.....	26
4.3.3 Resources.....	15	5.5 Monitoring Jobs.....	27
4.4 Trajectory Analysis.....	16	5.6 Reading Logs.....	29
4.4.1 Input.....	16	5.7 Killing Jobs.....	30
4.4.2 Output.....	16	5.7.1 Kill an individual job.....	30
4.4.3 Resources.....	16	5.7.2 Kill all jobs.....	30
4.5 Trajectory Interpolation.....	17	5.7.3 Kill jobs running a process.....	31
4.5.1 Input.....	17	5.8 Disabling Compute Resources.....	31
4.5.2 Output.....	17	A Troubleshooting.....	32
4.5.3 Resources.....	17	A.1 Unscheduleable.....	32
4.6 Sector Intersections.....	17	A.2 OOMKilled.....	32
4.6.1 Input.....	18	A.3 CrashLoopBackoff.....	32
4.6.2 Output.....	18		

Table of Figures

Figure 1: Trajectories Pipeline Processes.	6	Figure 11:	
Figure 2: GCP Kubernetes Clusters page.	8	find_sector_intersections_on_day.....	19
Figure 3: GCP Storage Browser.....	11	Figure 12: import_apds_file.....	20
Figure 4: import_data_on_day.....	12	Figure 13: merge_apds_data_on_day.....	21
Figure 5: match_overnight_flights_on_day		Figure 14: GCP Compute Node Pricing...	23
.....	14	Figure 15: A pre-emptible high-mem node	
Figure 6: merge_overnight_data_on_day	15	pool.....	25
Figure 7: analyse_positions_on_day.....	16	Figure 16: GCP Kubernetes Workloads	
Figure 8: interpolate_trajectory_on_day..	17	page.....	27
Figure 9: find_sector_intersections_on_day		Figure 17: GCP Kubernetes Engine Job	
.....	17	details page.....	29
Figure 10:		Figure 18: GCP Stackdriver log page.....	30
find_airport_intersections_on_day.....	18		

Revision History

Issue	Date	Description
Draft A	19 August 2018	Initial draft
Draft B	20 August 2018	Incorporate internal review comments.
Draft C	20 September 2018	Update intersection processing guide.
Draft D	21 September 2018	Incorporate internal review comments.
1.0.0	11 December 2018	Raised to issue 1.0.0 for contract close-out.

References

Number	Document
1	Amendment No. 1 to Contract No. 17-110452-C Trajectories Production
2	https://kubernetes.io/
3	https://cloud.google.com/
4	https://www.docker.com/
5	https://cloud.google.com/storage/docs/gsutil

Glossary

Item	Description
CSV	Comma Separated Variable
GCP	Google Cloud Platform
JSON	Java Script Object Notation

Section

1

1 Introduction

1. This document describes how to use the Trajectories Production Pipeline for Amendment No. 1 to Contract No. 17-110452-C Trajectories Production via Cloud-Based Analytics for Performance Monitoring and Review, see [1].
2. The Trajectories Production Pipeline merges trajectory data from different sources then: analyses, interpolates and finds intersections with the merged trajectories, see Figure 1.

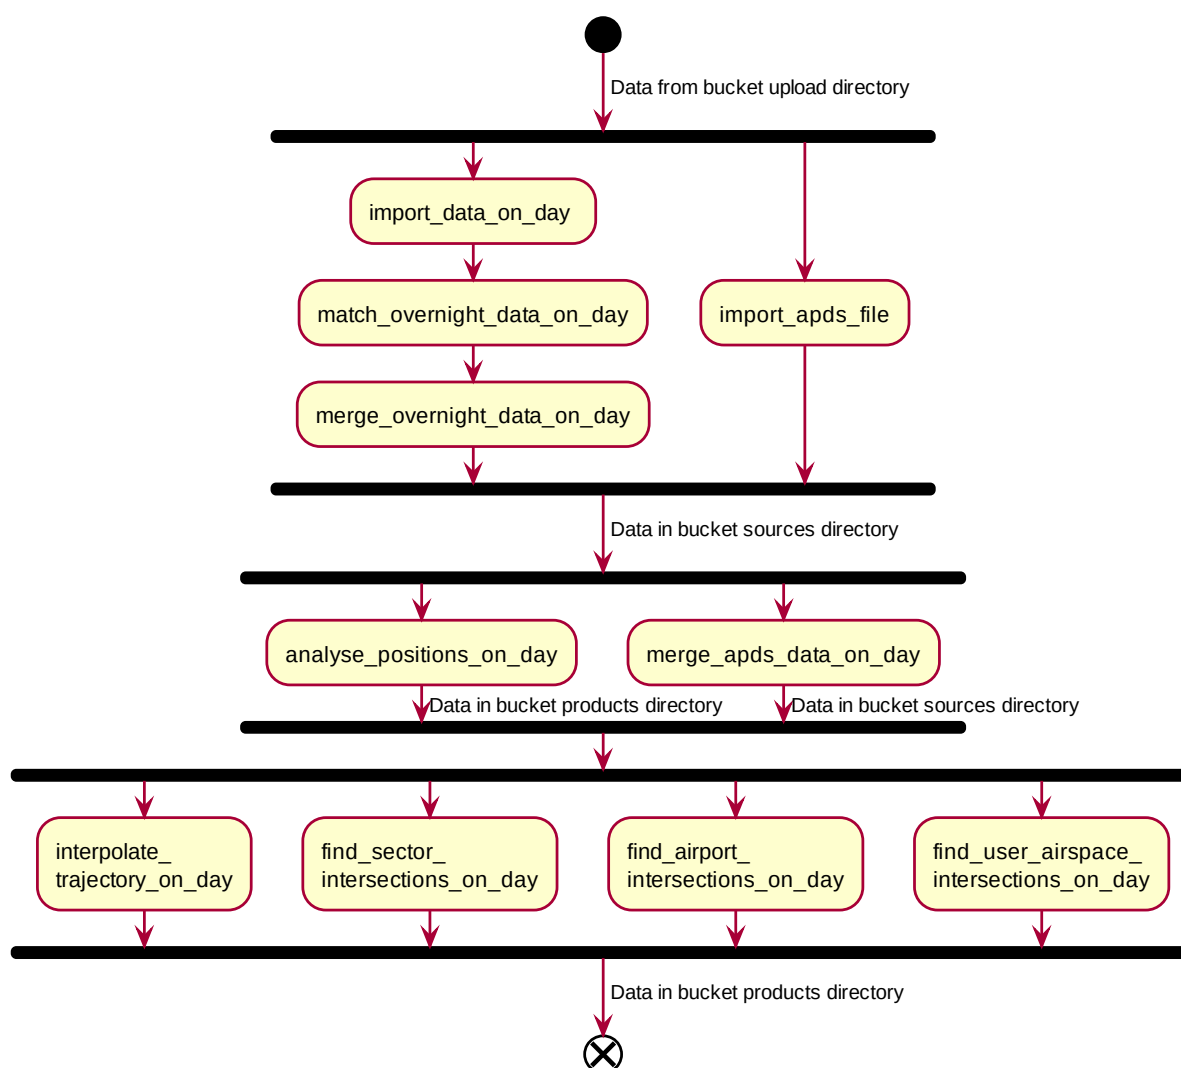


Figure 1: Trajectories Pipeline Processes

3. The Trajectories Production Pipeline runs on a Kubernetes [2] cluster on the Google Cloud Platform (GCP) [3].
4. A key feature of Kubernetes is its ability to support “horizontal scaling”, i.e. to run multiple processes in parallel.
5. The processes shown in Figure 1 must be run in the order shown for the same day’s data, e.g.: `match_overnight_data_on_day` must be run after `import_data_on_day` and before `merge_overnight_data_on_day`.
6. However, all the processes shown in Figure 1 can be run in parallel for different days data, e.g. `match_overnight_data_on_day` can be run in parallel for every day of a week, month, or even a year or more.

Section

2

2 System Setup

7. The Trajectories Production Pipeline runs on the Google Cloud Platform (GCP) [3].

2.1 GCP Project

8. GCP applications operate within the context of a project.

9. A project has a location which determines the geographical location of the associated project resources. The project must be named and a location must be selected.

10. Please note that a Google project has both a name an id and a number. The id may be auto-generated and will not necessarily be the same as the name. The name, id and number can be found on the gcloud console project information screen.

2.2 Google Bucket Storage

11. A Google bucket is used for long term data storage and backup storage.

12. This is associated with a project and must be created through the Gcloud console.

2.3 Kubernetes Cluster

13. The Trajectories Production Pipeline runs on a Kubernetes [2] cluster.

14. This is associated with a project and can be created through the Gcloud console, see Figure 2.

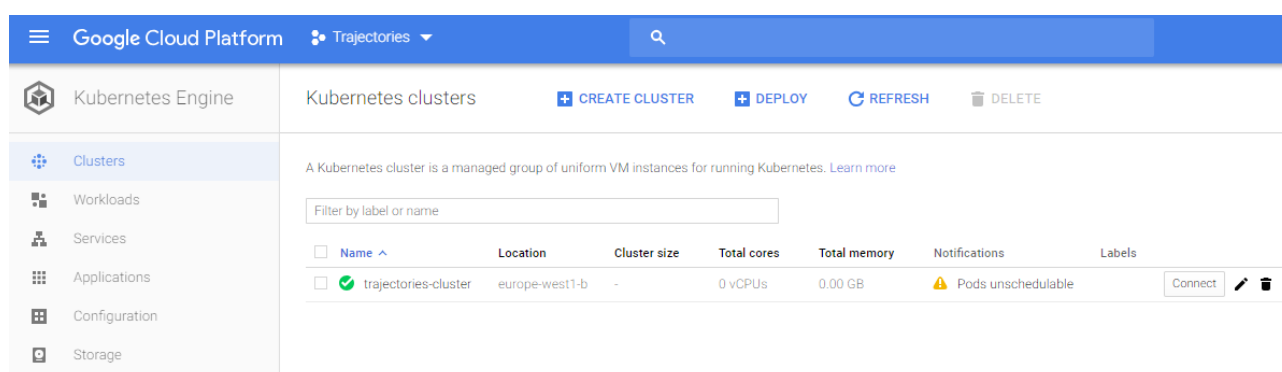


Figure 2: GCP Kubernetes Clusters page

2.4 Pipeline Container Image

15. The Trajectories Production Pipeline processes run on Kubernetes in a Docker [4] container. The docker container must be built and pushed to the relevant container registry for Kubernetes to be able to run the pipeline.
16. The Trajectories Production Pipeline runs a "python-only" container image. The container image is built in two parts:
 - the base image is built locally,
 - then the python code is built from the base image and pushed to the GCP container registry.

2.4.1 Build Base Image

17. The "python-only" base image is built from the BitBucket rt-config repository.
18. Clone the rt-config repository and from the rt-config directory, run:

```
docker build . -f docker/Dockerfile-python-base -t python-base:0.1.0
```

2.4.2 Build Python-only Image

19. The "python-only" image is built from the BitBucket rt-python repository.
20. Clone the rt-python repository and from the rt-python directory, run:

```
docker build . -f docker/Dockerfile-python-rt -t eu.gcr.io/<project>/python-rt:0.1.0
```

21. where <`project`> is the GCP project id, e.g.: shining-booth-205512.
22. To push the image to the GCP container registry, run:

```
docker push eu.gcr.io/<project>/python-rt:0.1.0
```

Section

3

3 Data Organisation

23. The Trajectories Production Pipeline processes read and write data to and from Google Bucket storage.

24. The bucket has the following directory structure:

```
pru-ta
|--airports
|  |--stands
|--airspaces
|  |--elementary
|  |--user_defined
|--products
|  |--error_metrics
|  |  |--cpr
|  |  |--cpr_fr24
|  |  |--|--overnight
|  |  |--fr24
|  |--fleet_data
|  |--intersections
|  |  |--airport
|  |  |  |--cpr
|  |  |  |--cpr_fr24
|  |  |  |--fr24
|  |  |--sector
|  |  |  |--cpr
|  |  |  |--cpr_fr24
|  |  |  |--fr24
|  |  |--user
|  |  |  |--cpr
|  |  |  |--cpr_fr24
|  |  |  |--fr24
|  |--synth_positions
|  |  |--cpr
|  |  |--cpr_fr24
|  |  |--fr24
|  |--traj_metrics
|  |  |--cpr
|  |  |--cpr_fr24
|  |  |--fr24
|  |--trajectories
|  |  |--cpr
|  |  |--cpr_fr24
|  |  |--fr24
|--refined
```

```

| | --apds
| | --cpr
| | --fr24
| | --merged
| | | --apds_cpr_fr24
| | | --daily_cpr_fr24
| | | --overnight_cpr_fr24
|--upload
| | --apds
| | --cpr
| | --fr24

```

25. Raw data must be copied into the relevant upload bucket before it is processed.

3.1 Bucket Storage Browser

26. The Google Buckets can be managed using the GCP storage browser, see Figure 3.

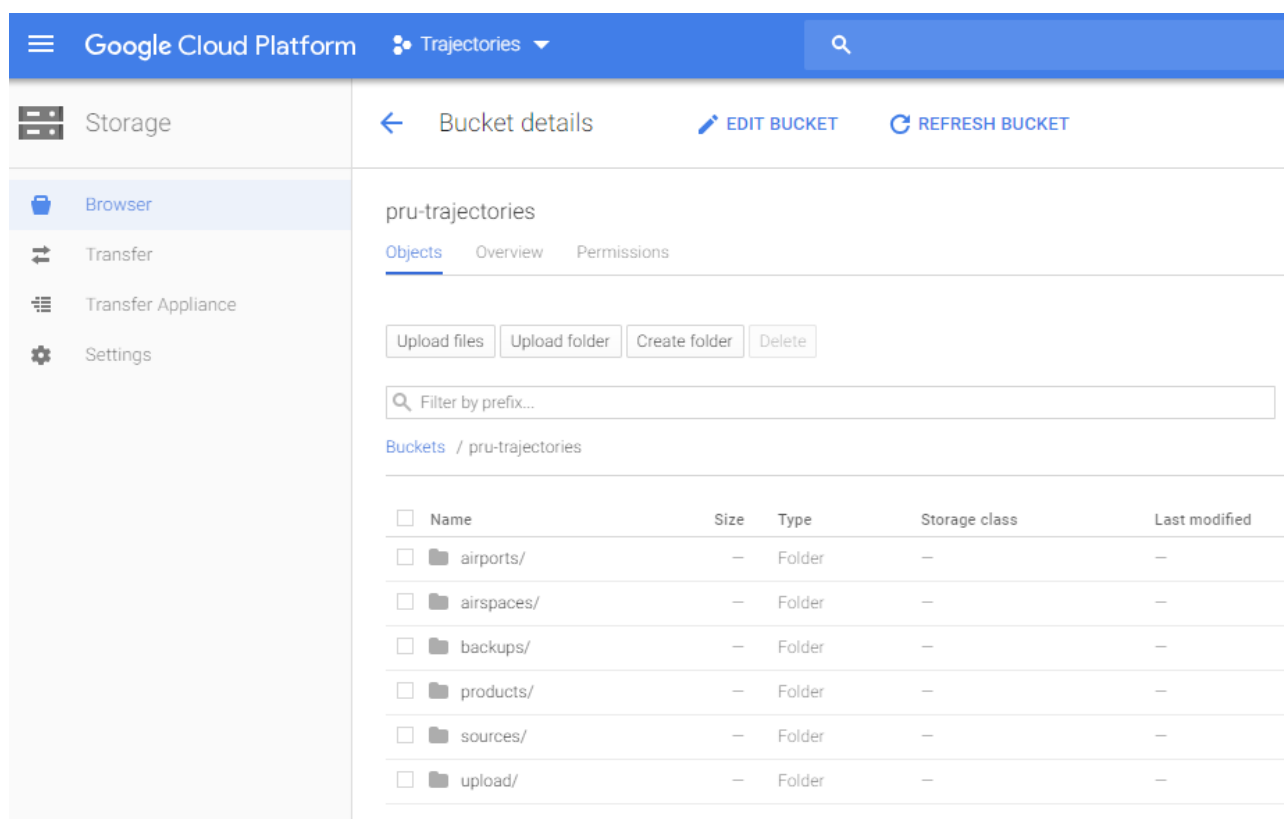


Figure 3: GCP Storage Browser

3.2 gsutil

27. Copying data to and from Google Buckets with the GCP storage browser is a relatively slow process for many large data files.

28. The gsutil tool [5] is a python application that lets you access GCP buckets from the command line, enabling multiple data files to be copied using scripts.

Section

4

4 Pipeline Processes

29. All the Trajectories Production Pipeline processes (see Figure 1) are run as kubernetes jobs.
30. However, most of the processes have different requirements to each other. I.e. some processes require files or even databases to be available, while some processes require significantly more computing resources (i.e. memory and/or cpu) than others.
31. The computing resources required by a process determines which GCP compute nodes the process can run on and how many processes kubernetes can run simultaneously as jobs. The less computing resources required by a process, the more jobs that can scaled horizontally by kubernetes, i.e. run in parallel.

4.1 Importing Daily Data

32. The `import_data_on_day` process refines and merges CPR and FR24 data, see Figure 4.

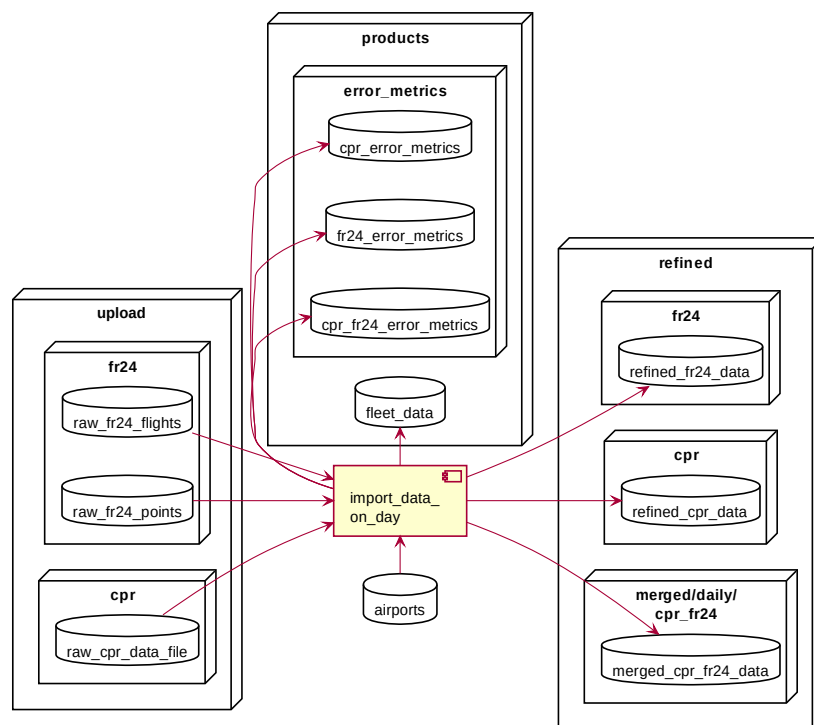


Figure 4: `import_data_on_day`

33. It runs the following applications:

- `convert_cpr_data.py`
- `convert_fr24_data.py`
- `convert_airport_ids.py`
- `extract_fleet_data.py`
- `clean_position_data.py`
- `match_cpr_adsb_trajectories.py`
- and `merge_cpr_adsb_trajectories.py`.

34. To refine and merge CPR and FR24 data for a given day.

4.1.1 Input

35. Data from the bucket upload directory:

- CPR data from `upload/cpr`
- and FR24 data from `upload/fr24`

36. Airport data from the airports bucket. The default file is `airports.csv`.

4.1.2 Output

37. Refined data in the bucket sources directory:

- CPR flights, positions and events files in the bucket `sources/cpr` directory,
- FR24 flights and positions files in the bucket `refined/fr24` directory,
- Merged CPR and FR24 flights, positions and events files in the bucket `refined/merged/daily_cpr_fr24` directory,

38. Product data in the bucket products directory:

- fleet data from the FR24 flights in the bucket `products/fleet_data`
- error metrics from cleaning the CPR, FR24 and merged CPR_FR24 positions files in the bucket `products/error_metrics` directory in the `cpr`, `fr24` and `cpr_fr24` directories respectively.

4.1.3 Resources

39. This process requires at least 10Gi memory and 1 cpu (note: it can use up to 3 cpus).

40. Each job normally runs to completion in around 1.5 hours.

4.2 Overnight Matching

41. The **match_overnight_flights_on_day** process matches flights for the given day with flights from the previous day and extracts flights that departed on the previous day to be merged later, see Figure 5.

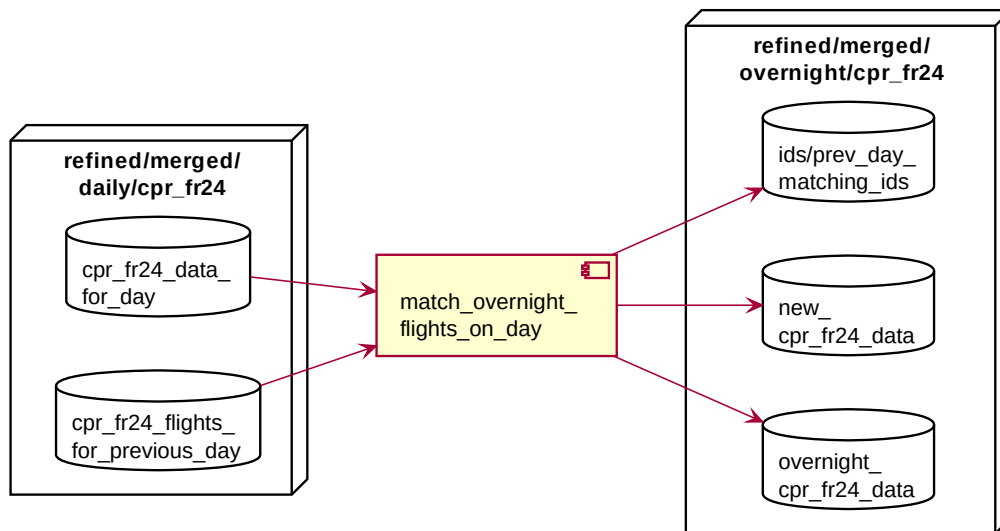


Figure 5: match_overnight_flights_on_day

42. It runs the **match_overnight_flights.py** and **extract_overnight_data.py** applications.

4.2.1 Inputs

43. Data from the bucket **refined/merged/daily_cpr_fr24** directory, i.e.: flights, positions and events

4.2.2 Output

44. Matched data in the bucket **refined/merged/overnight_cpr_fr24** directory:

- overnight positions and events for the previous day
- and “new” flights, positions and events for the day, without the flights for the previous day

45. Matching flight ids in the bucket **refined/merged/overnight_cpr_fr24/ids** directory.

4.2.3 Resources

46. This process requires at least 10Gi memory and 1 cpu.

47. Each job normally runs to completion in around 15 minutes.

4.3 Overnight Merging

48. The **merge_overnight_data_on_day** process merges data for the given day with data extracted from the next day by the **match_overnight_flights_on_day** process, see Figure 6.




Figure 6: merge_overnight_data_on_day

49. It runs the **merge_overnight_flight_data.py** and **clean_position_data.py** applications.

4.3.1 Input

50. Data in the bucket **refined/merged/overnight_cpr_fr24** directory:

- the “new” flights, positions and events for the day, without the flights for the previous day
- overnight positions and events from the next day.

51. Note: “new” flights, positions and events for a day are copied into the **sources/merged/overnight_cpr_fr24** directory by **match_overnight_flights_on_day**. Since **match_overnight_flights_on_day** cannot be run for the first day, the “new” flights, positions and events for first day must be copied from the **sources/merged/daily_cpr_fr24** directory with “new_” prepended to their filenames.

4.3.2 Output

52. Merged data for the day in the bucket **refined/merged/overnight_cpr_fr24** directory without “new_” or “overnight_” prepended to their filenames.

4.3.3 Resources

53. This process requires at least 10Gi memory and 1 cpu.

54. Each job normally runs to completion in around 40 minutes.

4.4 Trajectory Analysis

55. The **analyse_positions_on_day** process analyses position data for the given day, see Figure 7.

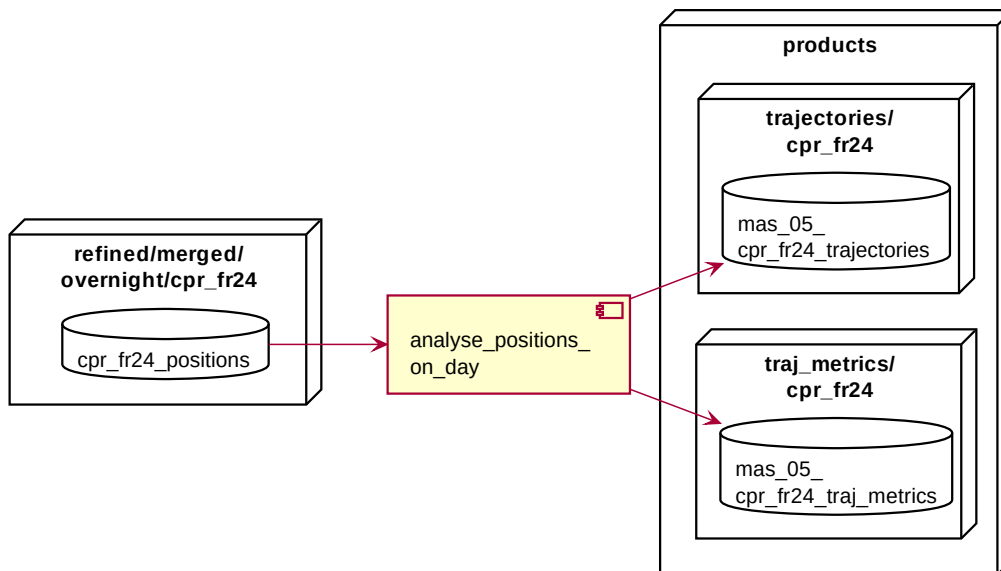


Figure 7: analyse_positions_on_day

56. It runs the analyse_position_data.py application.

4.4.1 Input

57. Merged overnight CPR/FR24 data for the day in the bucket refined/merged/overnight_cpr_fr24.

4.4.2 Output

58. A JSON trajectories file for the day in the bucket products/trajectories/cpr_fr24.

59. A CSV trajectory metrics file in the bucket products/traj_metrics/cpr_fr24.

4.4.3 Resources

60. This process only requires 500Mi memory and up to 1 cpu.

61. Each job normally runs to completion in around 30 minutes.

4.5 Trajectory Interpolation

62. The **interpolate_trajectory_on_day** process interpolates trajectory data for the given day, see Figure 8.

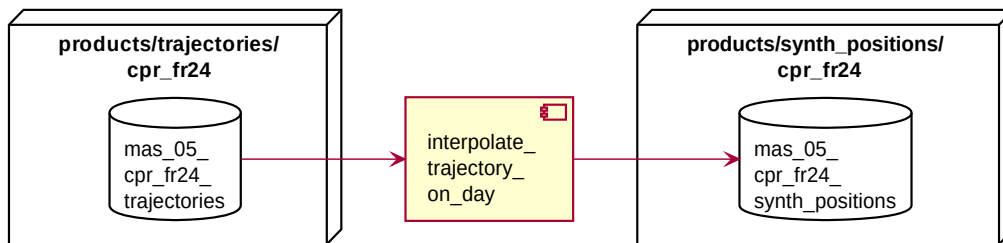


Figure 8: interpolate_trajectory_on_day

63. It runs the `interpolate_trajectories.py` application.

4.5.1 Input

64. The JSON trajectories file for the day from the bucket `products/trajectories/cpr_fr24`.

4.5.2 Output

65. A CSV synthetic positions file in the bucket `products/synth_positions/cpr_fr24`.

4.5.3 Resources

66. This process only requires 500Mi memory and up to 1 cpu.

67. Each job normally runs to completion in around 2 hours.

4.6 Sector Intersections

68. The **find_sector_intersections_on_day** process finds sector intersections for trajectory data for the given day, see Figure 9.

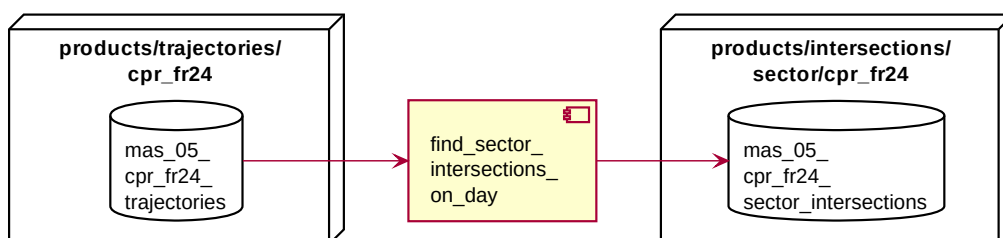


Figure 9: find_sector_intersections_on_day

69. It runs the `find_sector_intersections.py` application.

4.6.1 Input

70. The JSON trajectories file for the day from the bucket products/trajectories/cpr_fr24.

71. The airspace database must be running and loaded with the airspace sectors.

4.6.2 Output

72. A CSV sector intersections file in the bucket products/intersections/sector/cpr_fr24.

4.6.3 Resources

73. This process only requires 500Mi memory and up to 1 cpu.

74. The process requires the airspace database to be running, ideally on a compute node with at least one cpu core for each days data, e.g.: an n1-highcpu-32 compute node if being run for an entire AIRAC cycle. Note: the processes can run on the same compute node as the airspace database.

75. Each job normally runs to completion in around 18 hours.

4.7 Airport Intersections

76. The **find_airport_intersections_on_day** process finds airport intersections for trajectory data for the given day, see Figure 10.




Figure 10: find_airport_intersections_on_day

77. It runs the find_airport_intersections.py application.

4.7.1 Input

78. The JSON trajectories file for the day from the bucket products/trajectories/cpr_fr24.

4.7.2 Output

79. A CSV airport intersections file in the bucket products/intersections/airport/cpr_fr24.

4.7.3 Resources

80. This process only requires 500Mi memory and up to 1 cpu.

81. Each job normally runs to completion in around 15 minutes.

4.8 User Airspace Intersections

82. The `find_user_airspace_intersections_on_day` process finds user defined intersections for trajectory data for the given day, see Figure 11.

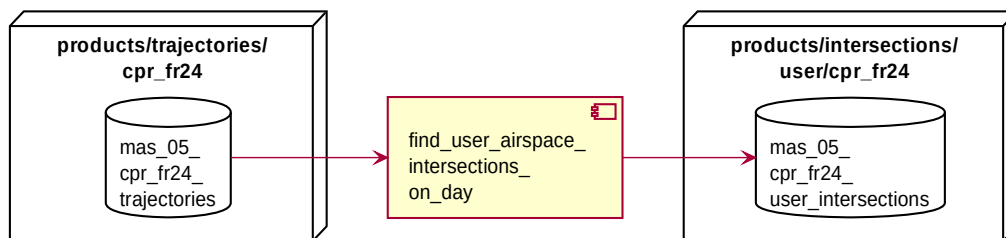


Figure 11: `find_sector_intersections_on_day`

83. It runs the `find_user_airspace_intersections.py` application.

4.8.1 Input

84. The JSON trajectories file for the day from the bucket `products/trajectories/cpr_fr24`.

85. The airspace database must be running and loaded with the user defined airspaces.

4.8.2 Output

86. A CSV user airspace intersections file in the bucket `products/intersections/user/cpr_fr24`.

4.8.3 Resources

87. This process only requires 500Mi memory and up to 1 cpu.

88. The process requires the airspace database to be running, ideally on a compute node with at least one cpu core for each days data, e.g.: an `n1-highcpu-32` machine if being run for an entire AIRAC cycle. Note: the processes can run on the same machine as the airspace database.

89. The time take for a job to run to completion depends upon the complexity of the airspace.

4.9 Import APDS Data

90. The **import_apds_file** process imports APDS airport data, see Figure 12.




Figure 12: import_apds_file

91. It runs the `convert_apd_data.py` application.

4.9.1 Input

92. The raw APDS data file, e.g.: `FAC_APDS_FLIGHT_IR691_2017-08-01_2017-09-01.csv.bz2`.

93. Airport stands data from the `airports/stands` bucket. The default file is `stands_EGLL.csv`.

4.9.2 Output

94. APDS flight, event and positions files for the period of the input file, e.g.:

- `apds_flights_2017-08-01_2017-09-01.csv`
- `apds_events_2017-08-01_2017-09-01.csv`
- `apds_positions_2017-08-01_2017-09-01.csv`

4.9.3 Resources

95. This process requires at least 10Gi memory and 1 cpu.

96. Each job normally runs to completion in under 30 minutes.

4.10 Merge APDS Data

97. The **merge_apds_data_on_day** process merges APDS data with the merged

CPR/FR24 data for the given day, see Figure 13.

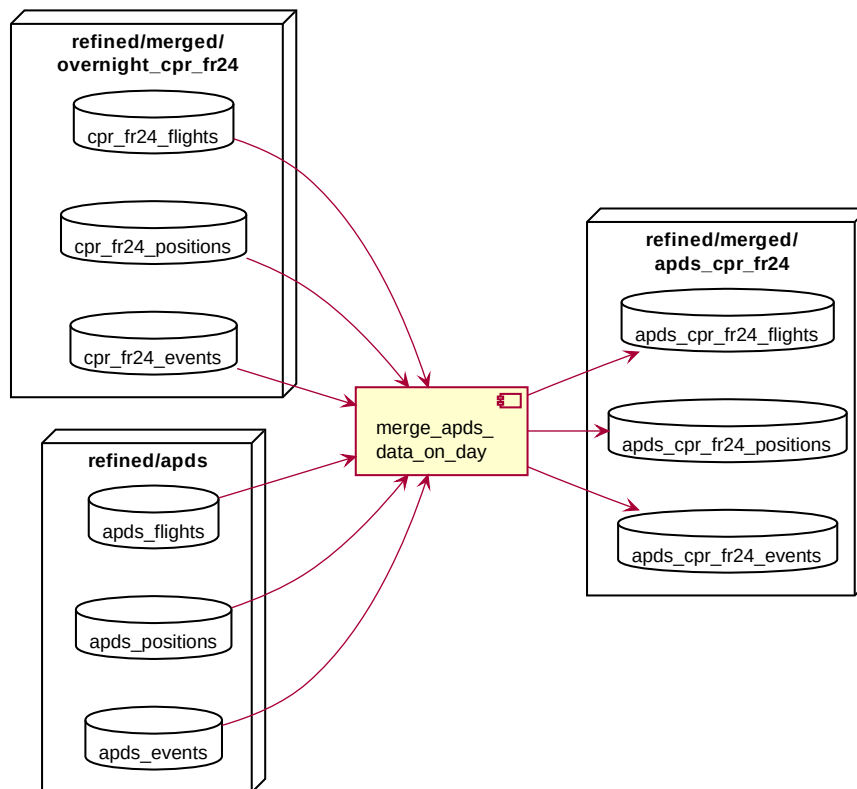


Figure 13: merge_apds_data_on_day

98. It runs the match_aptrajectories.py and merge_aptrajectories.py applications.

4.10.1 Input

99. APDS flight, event and positions files for the period of the input file, e.g.:

- apds_flights_2017-08-01_2017-09-01.csv
- apds_events_2017-08-01_2017-09-01.csv
- apds_positions_2017-08-01_2017-09-01.csv

100. Merged CPR and ADSB flight, event and positions files for the given dates.

4.10.2 Output

101. Merged APDS, CPR and ADSB flight, event and positions files for the given dates

4.10.3 Resources

102. This process requires at least 10Gi memory and 1 cpu.

103. Each job normally runs to completion in under 30 minutes.

Section

5

5 Managing Jobs

5.1 Processing Order

104. The CPR and FR24 data must be imported and processed in the order shown in Table 1.

Process	2017-09-01	2017-09-02	2017-09-03	2017-09-04	2017-09-05	2017-09-06	2017-09-07
import_data_on_day							
match_overnight_flights_on_day	See Note 1						
merge_overnight_data_on_day	See Note 2						See Note 3
analyse_positions_on_day							
interpolate_trajectory_on_day							

Table 1: Import Processing Order

105. Table 1 shows which processes can be run in parallel for a weeks data. However, processes may be run in parallel for as many days as input data is available.

106. Note 1: match_overnight_flights_on_day requires data for the date and the *previous* date. Therefore, it cannot be run for the first day in a set of data (unless the days data was imported previously).

107. Note 2: merge_overnight_data_on_day requires data for the date and the *next* date. Since match_overnight_flights_on_day copies data into the merged/overnight/cpr_fr24 bucket for the given day, the data files for the first day in a set of data must be copied over manually to merged/overnight/cpr_fr24 bucket with “new_” prepended to the file names.

108. Note 3: since merge_overnight_data_on_day requires data for the date and the next date, it cannot be run for the last day in a set of data.

5.2 Compute Resources

109. Some processes require more compute resources (memory and cpu) than others. In Table 1, processes with high memory requirements are shown in red, while processes with low memory requirements are shown in green.

110. The high memory processes require at least 10 GB memory (they may require even more memory for short periods). The smallest standard GCP machine type that they can run on is an n1-standard-4 machine with 15GB of memory and 4 cpus.

111. The other processes are limited to 0.5GB and can run on any GCP machine type, including an f1-micro with only 600MB of memory and 0.2 cpu.

112. GCP compute nodes can be run as preemptible nodes: short-lived compute instances that may be shut down with only 30 seconds notice.
113. Preemptible compute nodes are up to 80% cheaper than normal compute nodes and are ideal for running batch jobs like the pipeline processes, see Figure 14.




Figure 14: GCP Compute Node Pricing

114. Note: although preemptible f1-micro compute nodes are the cheapest, they are not necessarily the best value, since they only provide 0.2 cpu. For most processes that can run on any GCP machine type, preemptible n1-standard-1 compute nodes theoretically provide the best value. However, despite claiming 1 cpu they only support 0.5cpu...
115. In practice, preemptible g1-small compute nodes or multiple jobs running on a preemptible highcpu compute node provide the best value.

5.2.1 High Memory Processes

116. The processes:
- import_data_on_day,
 - match_overnight_flights_on_day,
 - merge_overnight_data_on_day,
 - import_apds_file
 - and merge_apds_data_on_day
117. Each require a compute unit with at least 10Gi memory. The minimum GCP machine type that they can run on as a compute node is n1-highmem-2.
118. However, we recommend running pairs of high memory pipeline processes on n1-highmem-4 compute nodes as the processes can share the 26GB of memory available.

5.2.2 Low Memory Processes

119. The processes:
- `analyse_positions_on_day`,
 - `interpolate_trajectory_on_day`,
 - `find_airport_intersections_on_day`,
 - `find_sector_intersections_on_day`
 - and `find_user_airspace_intersections_on_day`.
120. Can be limited to 500Mi memory per day. The minimum GCP machine type that they can run on as a compute node is f1-micro. However, the f1-micro normally only receives 0.2 cpu, with a full cpu only allowed in short bursts.
121. Preemptible n1-standard-1 compute nodes provide better value for processes that run at 100% cpu, i.e.: `analyse_positions_on_day` and `interpolate_trajectory_on_day`.

5.2.3 Airspace Database

122. The processes:
- `find_sector_intersections_on_day`
 - and `find_user_airspace_intersections_on_day`.
123. Use the GIS Postgres airspace database to find intersections. Therefore, an instance of the airspace database must be running and loaded with the airspace data for the AIRAC cycle.
124. The airspace database is the performance bottleneck when running these processes. Therefore it is recommended to run it on a “highcpu” node type with at least as many virtual cpus as the number of days being processed.
125. Note: the processes can run on the same compute node as the airspace database. Therefore, the airspace database and `find_sector_intersections_on_day` can be run for all 28 days of an AIRAC cycle on a preemptible n1-highcpu-32 compute node.

5.3 Running Processes

5.3.1 Provisioning Compute Resources

126. Before running a process or a batch of processes, the kubernetes cluster must be provisioned with enough compute resources (i.e. nodes) to run the processes.
127. The kubernetes master should not run on a preemptible node. A single f1-micro compute node is normally sufficient to run the kubernetes master.
128. Each high memory process requires at least a n1-highmem-2 compute node, or half a n1-highmem-4, a quarter of a n1-highmem-8, etc.
129. Figure 15 shows the configuration of a pool of 12 preemptible n1-highmem-4 compute nodes.

pre-rt-node-pool-high-mem (12 nodes, version 1.8.12-gk... ^

Name	pre-rt-node-pool-high-mem		
Size	12		
Node version	1.8.12-gke.3	Change	
Node image	Container-Optimised OS (cos)	Change	
Machine type	n1-highmem-4 (4 vCPUs, 26 GB memory)		
Total cores	48 vCPUs		
Total memory	312.00 GB		
Automatic node upgrades	Disabled		
Automatic node repair	Enabled		
Auto-scaling	Off		
Pre-emptible nodes	Enabled		
Boot disk type	Standard persistent disk		
Boot disk size in GB (per node)	100		
Local SSD disks (per node)	0		
Instance groups	gke-int-ta-cluster-pre-rt-node-pool-h-411cbae5-grp		

Kubernetes labels
No labels set

Taints
No taints set

Figure 15: A pre-emptible high-mem node pool

130. Each low memory process can run on a preemptible f1-micro compute node. They will run quite happily on high memory compute nodes, but they are much more expensive, see Figure 14.
131. The exception being the airspace database processes (see section 5.2.3), which require a compute node with enough virtual cpus for each day to be processed.
132. GCP takes around a minute to shutdown each compute node in a node pool: i.e. it is much quicker to shut down a single highcpu compute node than ran multiple jobs than many small or micro compute nodes than ran single jobs. However, some jobs such as `interpolate_trajectory_on_day` must be run on multiple compute nodes.

5.4 Running Jobs

133. Pipeline processes can be run on the kubernetes compute nodes using the `run_kubernetes_jobs.py` python script in the `rt-config` repo scripts directory. It is used as follows:

```
run_kubernetes_jobs <process name> <start date> <finish date>
```

134. where process name is the name of the process: e.g. `import_data_on_day` and start date and finish date are the first and last dates (in ISO 8601 format) of the data to run the process on.

135. For example:

```
run_kubernetes_jobs.py import_data_on_day 2017-08-07 2017-08-31
```

136. Runs the `import_data_on_day` process on data for all days between 2017-08-07 and 2017-08-31 inclusive and outputs:

```
('job "pru-import-data-on-day-2017-08-07-job-z8mqt" created\n', None)
('job "pru-import-data-on-day-2017-08-08-job-pzgvt" created\n', None)
('job "pru-import-data-on-day-2017-08-09-job-v8hqj" created\n', None)
('job "pru-import-data-on-day-2017-08-10-job-8vwpz" created\n', None)
('job "pru-import-data-on-day-2017-08-11-job-qprzt" created\n', None)
('job "pru-import-data-on-day-2017-08-12-job-gchvv" created\n', None)
('job "pru-import-data-on-day-2017-08-13-job-6kp7f" created\n', None)
('job "pru-import-data-on-day-2017-08-14-job-f6dcr" created\n', None)
('job "pru-import-data-on-day-2017-08-15-job-qnwxx" created\n', None)
('job "pru-import-data-on-day-2017-08-16-job-rmp8f" created\n', None)
('job "pru-import-data-on-day-2017-08-17-job-9xwpp" created\n', None)
('job "pru-import-data-on-day-2017-08-18-job-vcv8h" created\n', None)
('job "pru-import-data-on-day-2017-08-19-job-hgvtj" created\n', None)
('job "pru-import-data-on-day-2017-08-20-job-9nmdj" created\n', None)
('job "pru-import-data-on-day-2017-08-21-job-qvrcg" created\n', None)
('job "pru-import-data-on-day-2017-08-22-job-2jgrm" created\n', None)
('job "pru-import-data-on-day-2017-08-23-job-jwwk5" created\n', None)
('job "pru-import-data-on-day-2017-08-24-job-7z992" created\n', None)
('job "pru-import-data-on-day-2017-08-25-job-p629k" created\n', None)
('job "pru-import-data-on-day-2017-08-26-job-r4pf4" created\n', None)
('job "pru-import-data-on-day-2017-08-27-job-sr6g4" created\n', None)
('job "pru-import-data-on-day-2017-08-28-job-d6h64" created\n', None)
('job "pru-import-data-on-day-2017-08-29-job-nsdtd" created\n', None)
('job "pru-import-data-on-day-2017-08-30-job-vfq9g" created\n', None)
('job "pru-import-data-on-day-2017-08-31-job-dh8z8" created\n', None)
```

137. Which lists all of the created kubernetes jobs.

138. Note: `import_apds_file` takes the name of the APDS file, e.g.: `FAC_APDS_FLIGHT_IR691_2017-08-01_2017-09-01.csv.bz2`

139. Also note: where `merge_apds_data_on_day` takes two dates, they must be the start and end dates from the filename given to `import_apds_file`. `merge_apds_data_on_day` can also take four dates, where the first two dates are the date range and the last two dates are from the apds filename.

5.5 Monitoring Jobs

140. The jobs can be monitored from the GCP Kubernetes Engine Workloads page, see Figure 16.
141. Figure 16 shows a number of successfully launched jobs that are still running.
142. Whether a job launched successfully or not can be determined from the “status” column. Appendix A Troubleshooting describes some common errors and how they can be resolved.
143. Whether a job is still running or not can be determined from the “pods” column: “1/1” indicates that the job is still running, “0/1” indicates that the job has completed successfully.

Kubernetes Engine

Clusters

Workloads

Services

Applications

Configuration

Storage

Marketplace

Workloads

REFRESH

DEPLOY

Workloads are deployable units of computing that can be created and managed in a cluster.

Is system object : False

Filter workloads

Columns

Name	Status	Type	Pods	Namespace	Cluster
airspace-db	OK	Deployment	1/1	int	int-ta-cluster
hub	OK	Deployment	1/1	int	int-ta-cluster
int-ta-kube-kube-lego	OK	Deployment	1/1	int	int-ta-cluster
proxy	OK	Deployment	1/1	int	int-ta-cluster
pru-merge-overnight-data-on-day-2017-08-07-job-lwlwx	OK	Job	1/1	int	int-ta-cluster
pru-merge-overnight-data-on-day-2017-08-08-job-zhix9	OK	Job	1/1	int	int-ta-cluster
pru-merge-overnight-data-on-day-2017-08-09-job-6x6n4	OK	Job	1/1	int	int-ta-cluster
pru-merge-overnight-data-on-day-2017-08-10-job-g49pc	OK	Job	1/1	int	int-ta-cluster
pru-merge-overnight-data-on-day-2017-08-11-job-2vfft	OK	Job	1/1	int	int-ta-cluster
pru-merge-overnight-data-on-day-2017-08-12-job-jbthv	OK	Job	1/1	int	int-ta-cluster
pru-merge-overnight-data-on-day-2017-08-13-job-wcgp9	OK	Job	1/1	int	int-ta-cluster
pru-merge-overnight-data-on-day-2017-08-14-job-bkbrj	OK	Job	1/1	int	int-ta-cluster
pru-merge-overnight-data-on-day-2017-08-15-job-5v9zh	OK	Job	1/1	int	int-ta-cluster
pru-merge-overnight-data-on-day-2017-08-16-job-9wqt4	OK	Job	1/1	int	int-ta-cluster
pru-merge-overnight-data-on-day-2017-08-17-job-fjjnz	OK	Job	1/1	int	int-ta-cluster
pru-merge-overnight-data-on-day-2017-08-18-job-bl2ks	OK	Job	1/1	int	int-ta-cluster
pru-merge-overnight-data-on-day-2017-08-19-job-cz8gd	OK	Job	1/1	int	int-ta-cluster

Figure 16: GCP Kubernetes Workloads page

144. The status of running jobs can also be found from the command line by running:
- ```
kubectl get jobs -n <namespace> -L proc,date
```
145. which outputs:

| NAME<br>DATE                                | DESIRED | SUCCESSFUL | AGE | PROC |
|---------------------------------------------|---------|------------|-----|------|
| pru-import-data-on-day-2017-08-07-job-z8mqt | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-07               |         |            |     |      |
| pru-import-data-on-day-2017-08-08-job-pzgv  | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-08               |         |            |     |      |
| pru-import-data-on-day-2017-08-09-job-v8hqj | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-09               |         |            |     |      |
| pru-import-data-on-day-2017-08-10-job-8vwpz | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-10               |         |            |     |      |
| pru-import-data-on-day-2017-08-11-job-qprzt | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-11               |         |            |     |      |
| pru-import-data-on-day-2017-08-12-job-gchvv | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-12               |         |            |     |      |
| pru-import-data-on-day-2017-08-13-job-6kp7f | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-13               |         |            |     |      |
| pru-import-data-on-day-2017-08-14-job-f6dcr | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-14               |         |            |     |      |
| pru-import-data-on-day-2017-08-15-job-qnwxh | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-15               |         |            |     |      |
| pru-import-data-on-day-2017-08-16-job-rmp8f | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-16               |         |            |     |      |
| pru-import-data-on-day-2017-08-17-job-9xwpp | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-17               |         |            |     |      |
| pru-import-data-on-day-2017-08-18-job-vcv8h | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-18               |         |            |     |      |
| pru-import-data-on-day-2017-08-19-job-hgvtj | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-19               |         |            |     |      |
| pru-import-data-on-day-2017-08-20-job-9nmdj | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-20               |         |            |     |      |
| pru-import-data-on-day-2017-08-21-job-qvrcg | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-21               |         |            |     |      |
| pru-import-data-on-day-2017-08-22-job-2jgrm | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-22               |         |            |     |      |
| pru-import-data-on-day-2017-08-23-job-jwwk5 | 1       | 0          | 1h  |      |
| import_data_on_day 2017-08-23               |         |            |     |      |
| pru-import-data-on-day-2017-08-24-job-7z992 | 1       | 0          | 1h  |      |
| import_data_on_day 2017-08-24               |         |            |     |      |
| pru-import-data-on-day-2017-08-25-job-p629k | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-25               |         |            |     |      |
| pru-import-data-on-day-2017-08-26-job-r4pf4 | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-26               |         |            |     |      |
| pru-import-data-on-day-2017-08-27-job-sr6g4 | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-27               |         |            |     |      |
| pru-import-data-on-day-2017-08-28-job-d6h64 | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-28               |         |            |     |      |
| pru-import-data-on-day-2017-08-29-job-nsdtd | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-29               |         |            |     |      |
| pru-import-data-on-day-2017-08-30-job-vfq9g | 1       | 1          | 1h  |      |
| import_data_on_day 2017-08-30               |         |            |     |      |
| pru-import-data-on-day-2017-08-31-job-dh8z8 | 1       | 0          | 1h  |      |
| import_data_on_day 2017-08-31               |         |            |     |      |

146. Where a completed job has a 1 in the “SUCCESSFUL” column.

## 5.6 Reading Logs

147. Kubernetes logs the console output of all pods and jobs.
148. The log for a job can be read by selecting the job name on the GCP Kubernetes Engine Workloads page (see Figure 16).
149. This brings up the Job details page for the selected job, see Figure 17.

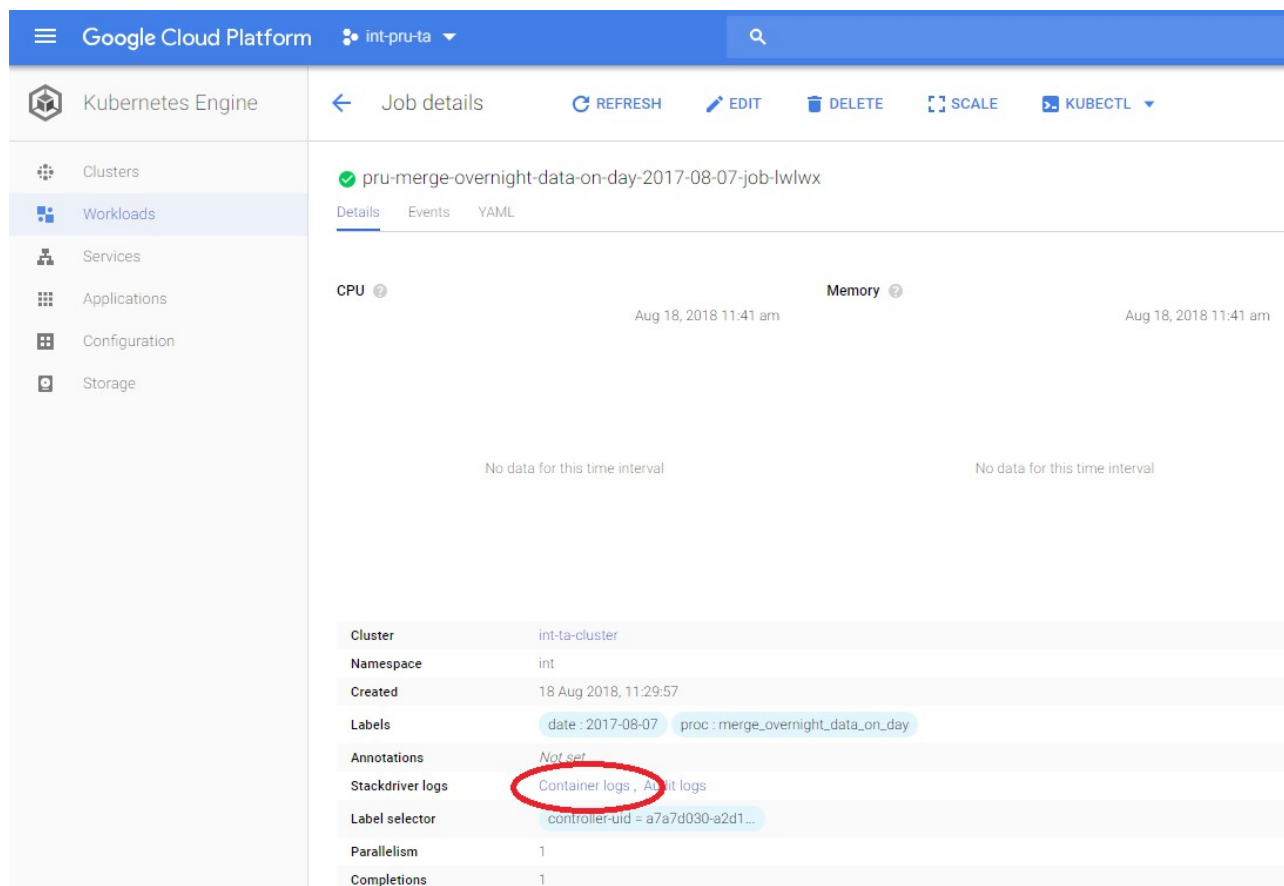


Figure 17: GCP Kubernetes Engine Job details page

150. Clicking on “Container logs” on the Job details page (see red ellipse in Figure 17) brings up the GCP Stackdriver Logging page, see Figure 18.
151. The log usually opens at the most recent entry. The down arrow by the “Jump to now” button can be used to “Jump to first entry” and the scroll bar can be used to scroll through log entries.

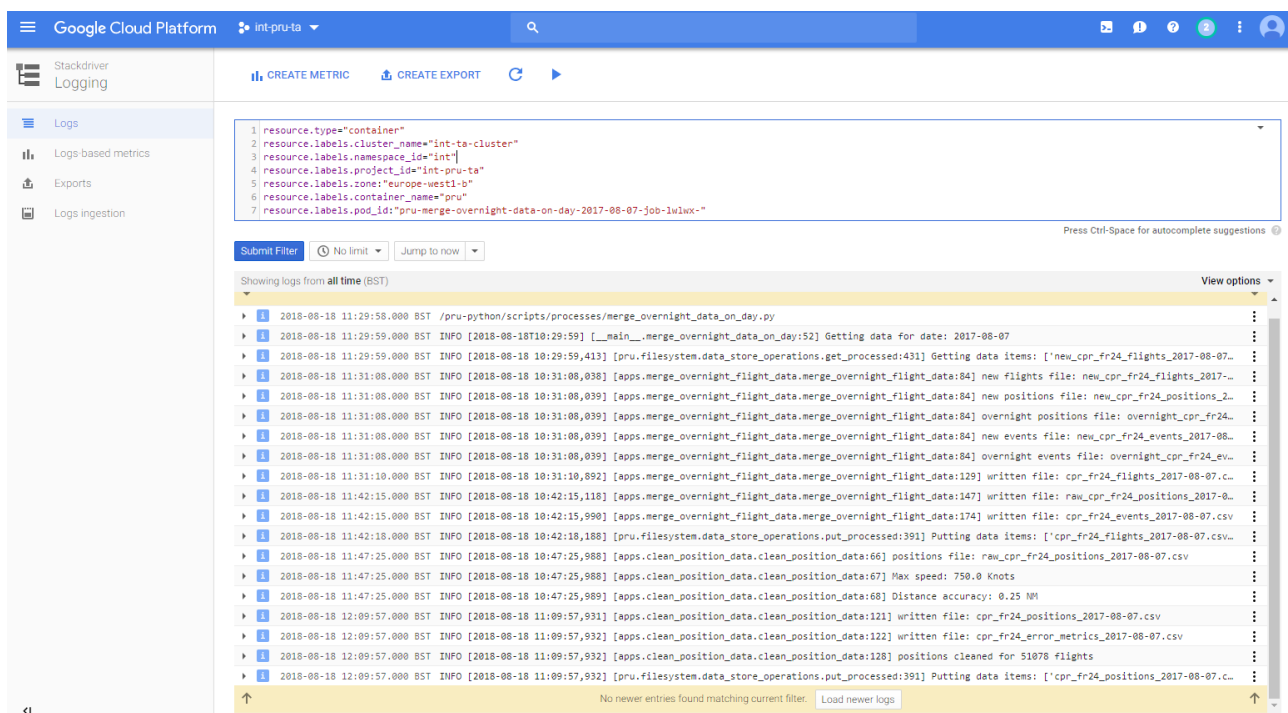


Figure 18: GCP Stackdriver log page

## 5.7 Killing Jobs

152. Kubernetes does not kill completed jobs. The jobs are left so that there logs can be inspected. The user must kill completed jobs manually.

### 5.7.1 Kill an individual job

153. An individual job can be killed with the following command:

```
kubectl delete jobs -n <namespace> <job name>
```

154. For example:

```
kubectl delete jobs -n int pru-import-data-on-day-2017-08-07-job-z8mqt
```

### 5.7.2 Kill all jobs

155. All jobs can be killed with the following command:

```
kubectl delete jobs -n <namespace> --all
```

156. For example:

```
kubectl delete jobs -n int --all
```

### 5.7.3 Kill jobs running a process

157. All the jobs running a process can be killed with:

```
kubect1 delete jobs -n <namespace> -l proc=<process name>
```

158. For example:

```
kubect1 delete jobs -n int -l proc=import_data_on_day
```

## 5.8 Disabling Compute Resources

159. GCP charges for all the compute nodes enabled in a Kubernetes cluster regardless of whether they are used or not. Therefore, it is recommended to disable unused compute nodes, by editing the cluster configuration, setting their node pool sizes to zero, see Figure 15

160. Note: GCP takes around a minute to shutdown each node in a node pool, e.g. it may take over 30 minutes to shutdown low memory process nodes for a months data if they use individual compute nodes for each kubernetes job.

**Appendix****A**

# **A Troubleshooting**

161. The following are some common errors encountered while running jobs on kubernetes and what you can do to resolve them.

## **A.1 Unscheduleable**

162. This error is encountered when there are not enough resources for kubernetes to schedule a job.

163. Clicking on the word “unschedulable” should tell you why the job could not be scheduled, i.e. “insufficient cpu” or “insufficient memory”, or both.

### **A.1.1 Solution**

164. First wait a minute to ensure that the job really is unschedulable. Often kubernetes simply needs more time to schedule a job.

165. The solution is to provision more compute resources or kill the job and run it again after other jobs have finished.

## **A.2 OOMKilled**

166. This error is encountered when a running job exceeds its kubernetes memory resource limit.

### **A.2.1 Solution**

167. The job needs more memory to run, consider increasing kubernetes memory requests/and or limits for the job.

## **A.3 CrashLoopBackoff**

168. This error is encountered when there is an issue with the process running in the container. For example a PostGIS database not being available or a bug the pipeline process.

### **A.3.1 Solution**

169. Collect logs from the job and contact Via Technology.